



```
In [12]: import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import sqlite3
import csv
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
#from wordcloud import WordCloud
import re
import os
from sqlalchemy import create_engine # database connection
import datetime as dt
import joblib
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import SGDClassifier
from sklearn import metrics
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from skmultilearn.adapt import mlknn
from skmultilearn.problem_transform import ClassifierChain
from skmultilearn.problem_transform import BinaryRelevance
from skmultilearn.problem_transform import LabelPowerset
from sklearn.naive_bayes import GaussianNB
from datetime import datetime
```

# Stack Overflow: Tag Prediction

## 1. Business Problem

### 1.1 Description

#### Description

Stack Overflow is the largest, most trusted online community for developers to learn, share their programming knowledge, and build their careers.

Stack Overflow is something which every programmer use one way or another. Each month, over 50 million developers come to Stack Overflow to learn, share their knowledge, and build their careers. It features questions and answers on a wide range of topics in computer programming.

The website serves as a platform for users to ask and answer questions, and, through membership and active participation, to vote questions and answers up or down and edit questions and answers in a fashion similar to a wiki or Digg. As of April 2014 Stack Overflow has over 4,000,000 registered users, and it exceeded 10,000,000 questions in late August 2015. Based on the type of tags assigned to questions, the top eight most discussed topics on the site are: Java, JavaScript, C#, PHP, Android, jQuery, Python and HTML.

## Problem Statemtent

Suggest the tags based on the content that was there in the question posted in Stackoverflow.

**Source:** <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/>

## 1.2 Sources / useful links

Data Source : <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>

(<https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>)

youtube : <https://youtu.be/nNDqbUhtIRg> (<https://youtu.be/nNDqbUhtIRg>)

research paper : <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf> (<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>)

research paper : <https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL> (<https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL>)

## 1.3 Real World / Business Objectives and Constraints

1. Predict as man labels as possible correctly.
2. No strict latency constaraint.
3. Cost of errors would be a bad customer experience.

## 2. Machine Learning problem

### 2.1 Data

#### 2.1.1 Data Overview

All of the data is in 2 files: Train and Test.

**Train.csv** contains 4 columns: Id,Title,Body,Tags.

**Test.csv** contains the same columns but without the Tags, which you are to predict.

**Size of Train.csv** - 6.75GB

**Size of Test.csv** - 2GB

**Number of rows in Train.csv** = 6034195

The questions are randomized and contains a mix of verbose text sites as well as sites related to math and programming. The number of questions from each site may vary, and no filtering has been performed on the questions (such as closed questions).

### Data Field Explanation

Dataset contains 6034195 rows. The column in the table are:

**Id** - Unique identifier for each question

**Title** - The question's title

**Body** - The body of the question

**Tags** - The tags associated with the question (all lowercase, should not contain tabs '\t' or ampersands '&')

### 2.1.2 Example Data point

**Title:** Implementing Boundary Value Analysis of Software Testing in a C++ program?

**Body :**

```

#include<
iostream>\n
#include<
stdlib.h>\n\n
using namespace std;\n\n
int main()\n
{\n
    int n,a[n],x,c,u[n],m[n],e[n][4];\n
    cout<<"Enter the number of variables";\n
cin>>n;\n\n
    cout<<"Enter the Lower, and Upper Limits of the
variables";\n
    for(int y=1; y<n+1; y++)\n
    {\n
        cin>>m[y];\n
        cin>>u[y];\n
    }\n
    for(x=1; x<n+1; x++)\n
    {\n
        a[x] = (m[x] + u[x])/2;\n
    }\n
    c=(n*4)-4;\n
    for(int a1=1; a1<n+1; a1++)\n
    {\n\n
        e[a1][0] = m[a1];\n
        e[a1][1] = m[a1]+1;\n
        e[a1][2] = u[a1]-1;\n
        e[a1][3] = u[a1];\n
    }\n
    for(int i=1; i<n+1; i++)\n
    {\n
        for(int l=1; l<=i; l++)\n
        {\n
            if(l!=1)\n
            {\n
                cout<<a[l]<<"\\t";\n
            }\n
        }\n
        for(int j=0; j<4; j++)\n
        {\n
            cout<<e[i][j];\n
            for(int k=0; k<n-(i+1); k++)\n
            {\n
                cout<<a[k]<<"\\t";\n
            }\n
            cout<<"\\n";\n
        }\n
    }\n
}    \n\n

```

```

system("PAUSE");\n
return 0;    \n
}\n

```

\n\n

<p>The answer should come in the form of a table like</p>\n\n

<pre><code>

```

1          50          50\n
2          50          50\n
99         50          50\n
100        50          50\n
50         1           50\n
50         2           50\n
50         99          50\n
50         100         50\n
50         50          1\n
50         50          2\n
50         50          99\n
50         50          100\n

```

</code></pre>\n\n

<p>if the no of inputs is 3 and their ranges are\n

1,100\n

1,100\n

1,100\n

(could be varied too)</p>\n\n

<p>The output is not coming,can anyone correct the code or tell me what's wrong?</p>\n'

Tags : 'c++ c'

## 2.2 Mapping the real-world problem to a Machine Learning Problem

### 2.2.1 Type of Machine Learning Problem

It is a multilable classification problem

**Multilable Classification:** Multilabel classification assigns to each sample a set of target labels. This can be thought as predicting properties of a data-point that are not mutually exclusive, such as topics that are relevant for a document. A text might be about any of religion, politics, finance or education at the same time or none of these.

\_\_\_Credit\_\_\_: <http://scikit-learn.org/stable/modules/multiclass.html>

### 2.2.2 Performance metric

**Micro-Averaged F1-Score (Mean F Score)** : The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0. The relative contribution of precision and recall to the F1 score are equal. The formula for the F1 score is:

$$F1 = 2 * (precision * recall) / (precision + recall)$$

In the multi-class and multi-label case, this is the weighted average of the F1 score of each class.

'micro f1 score':

Calculate metrics globally by counting the total true positives, false negatives and false positives.

'macro f1 score':

Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

<https://www.kaggle.com/wiki/MeanFScore> (<https://www.kaggle.com/wiki/MeanFScore>)

[http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1\\_score.html](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html) ([http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1\\_score.html](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html))

**Hamming loss** : The Hamming loss is the fraction of labels that are incorrectly predicted.

<https://www.kaggle.com/wiki/HammingLoss> (<https://www.kaggle.com/wiki/HammingLoss>)

## 2.2.3 Machine Learning Objectives and Constraints

1. Minimize Micro avg F1 Score.
2. Try out multiple strategies for Multi-label classification.

# 3. Exploratory Data Analysis

## 3.1 Data Loading and Cleaning

### 3.1.1 Using Pandas with SQLite to Load the data

```
In [18]: #Creating db file from csv
if not os.path.isfile('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train.db'):
    start = datetime.now()
    disk_engine = create_engine('sqlite:///D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train.csv', chunksize=chunksize):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1
    print("Time taken to run this cell :", datetime.now() - start)
```

### 3.1.2 Counting the number of rows

```
In [19]: if os.path.isfile('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train.db'):
    start = datetime.now()
    con = sqlite3.connect('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train.db')
    num_rows = pd.read_sql_query("SELECT count(*) FROM data", con)
    #Always remember to close the database
    print("Number of rows in the database :", "\n", num_rows['count(*)'].values[0])
    con.close()
    print("Time taken to count the number of rows :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cell to create the database")
```

Number of rows in the database :

6034196

Time taken to count the number of rows : 0:01:41.385089

### 3.1.3 Checking for duplicates

```
In [20]: if os.path.isfile('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train.db'):
    start = datetime.now()
    con = sqlite3.connect('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train.db')
    df_no_dup = pd.read_sql_query('SELECT Title, Body, Tags, COUNT(*) as cnt_dup FROM data', con)
    con.close()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the first cell to create the database")
```

Time taken to run this cell : 0:36:12.827058

```
In [21]: df_no_dup.head()
# we can observe that there are duplicates
```

```
Out[21]:
```

	Title	Body	Tags	cnt_dup
0	Implementing Boundary Value Analysis of S...	<pre>\n#include<iosstream>\n#include<...</pre>	c++ c	1
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding	1
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding columns	1
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in <a href="http://sta...	jsp jstl	1
4	java.sql.SQLException: [Microsoft] [ODBC Dri...	<p>I use the following code</p>\n\n<pre>\n<code>...</code>	java jdbc	2

```
In [22]: print("number of duplicate questions :", num_rows['count(*)'].values[0] - df_no_du
number of duplicate questions : 1827881 ( 30.292038906260256 % )
```

```
In [23]: # number of times each question appeared in our database
df_no_dup.cnt_dup.value_counts()
```

```
Out[23]: 1    2656284
2    1272336
3     277575
4         90
5         25
6          5
Name: cnt_dup, dtype: int64
```



```
In [24]: # creating bool series True for NaN values
bool_series = pd.isnull(df_no_dup["Tags"])

# filtering data
# displaying data only with team = NaN
print("Number of rows with NaN values = ", len(df_no_dup[bool_series]))
df_no_dup[bool_series]
```

Number of rows with NaN values = 7

Out[24]:

		Title	Body	Tags	cnt_dup
777547		Do we really need NULL?	<div>&lt;blockquote&gt;\n &lt;p&gt;&lt;strong&gt;Possible Duplicate:...</div>	None	1
962680	Find all values that are not null and not in a...		<div>&lt;p&gt;I am running into a problem which results i...</div>	None	1
1126558		Handle NullObjects	<div>&lt;p&gt;I have done quite a bit of research on best...</div>	None	1
1256102		How do Germans call null	<div>&lt;p&gt;In german null means 0, so how do they call...</div>	None	1
2430668	Page cannot be null. Please ensure that this o...		<div>&lt;p&gt;I get this error when i remove dynamically ...</div>	None	1
3329908	What is the difference between NULL and "0"?		<div>&lt;p&gt;What is the difference from NULL and "0"?&lt;/...&gt;</div>	None	1
3551595	a bit of difference between null and space		<div>&lt;p&gt;I was just reading this quote&lt;/p&gt;\n\n&lt;block...</div>	None	2

```
In [25]: # Deleting rows with NaN values
df_no_dup = df_no_dup.drop([777547, 962680, 1126558, 1256102, 2430668, 3329908, 3551595])
```

```
In [26]: # creating bool series True for NaN values
bool_series = pd.isnull(df_no_dup["Tags"])

# filtering data
# displaying data only with team = NaN
print("Number of rows with NaN values = ", len(df_no_dup[bool_series]))
```

Number of rows with NaN values = 0

```
In [27]: df_no_dup['index_col'] = df_no_dup.index
df_no_dup = df_no_dup.reset_index()
df_no_dup[777543:777548]
```

Out[27]:

	index	Title	Body	Tags	cnt_dup	index_col
<b>777543</b>	777543	Do we really need Automapper?	<p>I was learning AutoMapper and understand it...	automapper	2	777543
<b>777544</b>	777544	Do we really need Continuous Integration if we...	<p>Each of our team members runs all kind of t...	continuous-integration	1	777544
<b>777545</b>	777545	Do we really need Hahn-Banach that much?	<p>In the texts on functional analysis I'm rea...	abstract-algebra functional-analysis normed-sp...	2	777545
<b>777546</b>	777546	Do we really need MVC user controls, when movi...	<p>We are migrating and asp.NET application to...	asp.net asp.net-mvc	1	777546
<b>777547</b>	777548	Do we really need VOLATILE keyword in C#?	<p>Here is the code that I was trying on my wo...	c# multithreading volatile	2	777548

```
In [28]: df_no_dup = df_no_dup.drop("index", axis=1)
df_no_dup = df_no_dup.drop("index_col", axis=1)
df_no_dup[777543:777548]
```

Out[28]:

	Title	Body	Tags	cnt_dup
<b>777543</b>	Do we really need Automapper?	<p>I was learning AutoMapper and understand it...	automapper	2
<b>777544</b>	Do we really need Continuous Integration if we...	<p>Each of our team members runs all kind of t...	continuous-integration	1
<b>777545</b>	Do we really need Hahn-Banach that much?	<p>In the texts on functional analysis I'm rea...	abstract-algebra functional-analysis normed-sp...	2
<b>777546</b>	Do we really need MVC user controls, when movi...	<p>We are migrating and asp.NET application to...	asp.net asp.net-mvc	1
<b>777547</b>	Do we really need VOLATILE keyword in C#?	<p>Here is the code that I was trying on my wo...	c# multithreading volatile	2



```
In [32]: #This method seems more appropriate to work with this much data.
#creating the connection with database file.
if os.path.isfile('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train_no
    start = datetime.now()
    con = sqlite3.connect('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/t
    tag_data = pd.read_sql_query("""SELECT Tags FROM no_dup_train""", con)
    #Always remember to close the database
    con.close()

    # Let's now drop unwanted column.
    tag_data.drop(tag_data.index[0], inplace=True)
    #Printing first 5 columns from our data frame
    tag_data.head()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cells to
```

Time taken to run this cell : 0:03:16.984086

## 3.2 Analysis of Tags

### 3.2.1 Total number of unique tags

```
In [33]: # Importing & Initializing the "CountVectorizer" object, which
#is scikit-learn's bag of words tool.

#by default 'split()' will tokenize each tag using space.
vectorizer = CountVectorizer(tokenizer = lambda x: x.split())
# fit_transform() does two functions: First, it fits the model
# and learns the vocabulary; second, it transforms our training data
# into feature vectors. The input to fit_transform should be a list of strings.
tag_dtm = vectorizer.fit_transform(tag_data['Tags'])
```

```
In [34]: print("Number of data points :", tag_dtm.shape[0])
print("Number of unique tags :", tag_dtm.shape[1])
```

Number of data points : 4206307  
Number of unique tags : 42048

```
In [35]: # 'get_feature_name()' gives us the vocabulary.
tags = vectorizer.get_feature_names()
# Lets look at the tags we have.
print("Some of the tages we have :", tags[:10])
```

Some of the tages we have : ['.a', '.app', '.asp.net-mvc', '.aspxauth', '.bash-profile', '.class-file', '.cs-file', '.doc', '.drv', '.ds-store']

### 3.2.3 Number of times a tag appeared

```
In [36]: # https://stackoverflow.com/questions/15115765/how-to-access-sparse-matrix-elements
# Lets now store the document term matrix in a dictionary.
freqs = tag_dtm.sum(axis=0).A1
result = dict(zip(tags, freqs))
```


```
In [37]: # Saving this dictionary to csv files.
if not os.path.isfile('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/tag_counts.csv'):
    with open('D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/tag_counts.csv', 'w') as csv_file:
        writer = csv.writer(csv_file)
        for key, value in result.items():
            writer.writerow([key, value])
tag_df = pd.read_csv("D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/tag_counts.csv")
tag_df.head()
```

Out[37]:


	Tags	Counts
0	.a	18
1	.app	37
2	.asp.net-mvc	1
3	.aspxauth	21
4	.bash-profile	138

```
In [38]: tag_df_sorted = tag_df.sort_values(['Counts'], ascending=False)
tag_counts = tag_df_sorted['Counts'].values
```

```
In [39]: plt.plot(tag_counts)
plt.title("Distribution of number of times tag appeared questions")
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
```




```
In [40]: plt.plot(tag_counts[0:10000])
plt.title('first 10k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:10000:25]), tag_counts[0:10000:25])
```



400	331505	44829	22429	17728	13364	11162	10029	9148	8054	7151
6466	5865	5370	4983	4526	4281	4144	3929	3750	3593	
3453	3299	3123	2986	2891	2738	2647	2527	2431	2331	
2259	2186	2097	2020	1959	1900	1828	1770	1723	1673	
1631	1574	1532	1479	1448	1406	1365	1328	1300	1266	
1245	1222	1197	1181	1158	1139	1121	1101	1076	1056	
1038	1023	1006	983	966	952	938	926	911	891	
882	869	856	841	830	816	804	789	779	770	
752	743	733	725	712	702	688	678	671	658	
650	643	634	627	616	607	598	589	583	577	
568	559	552	545	540	533	526	518	512	506	
500	495	490	485	480	477	469	465	457	450	
447	442	437	432	426	422	418	413	408	403	
398	393	388	385	381	378	374	370	367	365	
361	357	354	350	347	344	342	339	336	332	
330	326	323	319	315	312	309	307	304	301	
299	296	293	291	289	286	284	281	278	276	
275	272	270	268	265	262	260	258	256	254	
252	250	249	247	245	243	241	239	238	236	
234	233	232	230	228	226	224	222	220	219	
217	215	214	212	210	209	207	205	204	203	
201	200	199	198	196	194	193	192	191	189	
188	186	185	183	182	181	180	179	178	177	
175	174	172	171	170	169	168	167	166	165	
164	162	161	160	159	158	157	156	156	155	
154	153	152	151	150	149	149	148	147	146	
145	144	143	142	142	141	140	139	138	137	
137	136	135	134	134	133	132	131	130	130	
129	128	128	127	126	126	125	124	124	123	
123	122	122	121	120	120	119	118	118	117	
117	116	116	115	115	114	113	113	112	111	
111	110	109	109	108	108	107	106	106	106	


105	105	104	104	103	103	102	102	101	101
100	100	99	99	98	98	97	97	96	96
95	95	94	94	93	93	93	92	92	91
91	90	90	89	89	88	88	87	87	86
86	86	85	85	84	84	83	83	83	82
82	82	81	81	80	80	80	79	79	78
78	78	78	77	77	76	76	76	75	75
75	74	74	74	73	73	73	73	72	72]

```
In [41]: plt.plot(tag_counts[0:1000])
plt.title('first 1k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:1000:5]), tag_counts[0:1000:5])
```



200	[331505	221533	122769	95160	62023	44829	37170	31897	26925	24537
22429	21820	20957	19758	18905	17728	15533	15097	14884	13703	
13364	13157	12407	11658	11228	11162	10863	10600	10350	10224	
10029	9884	9719	9411	9252	9148	9040	8617	8361	8163	
8054	7867	7702	7564	7274	7151	7052	6847	6656	6553	
6466	6291	6183	6093	5971	5865	5760	5577	5490	5411	
5370	5283	5207	5107	5066	4983	4891	4785	4658	4549	
4526	4487	4429	4335	4310	4281	4239	4228	4195	4159	
4144	4088	4050	4002	3957	3929	3874	3849	3818	3797	
3750	3703	3685	3658	3615	3593	3564	3521	3505	3483	
3453	3427	3396	3363	3326	3299	3272	3232	3196	3168	
3123	3094	3073	3050	3012	2986	2983	2953	2934	2903	
2891	2844	2819	2784	2754	2738	2726	2708	2681	2669	
2647	2621	2604	2594	2556	2527	2510	2482	2460	2444	
2431	2409	2395	2380	2363	2331	2312	2297	2290	2281	
2259	2246	2222	2211	2198	2186	2162	2142	2132	2107	
2097	2078	2057	2045	2036	2020	2011	1994	1971	1965	
1959	1952	1940	1932	1912	1900	1879	1865	1855	1841	
1828	1821	1813	1801	1782	1770	1760	1747	1741	1734	
1723	1707	1697	1688	1683	1673	1665	1656	1646	1639]	

```
In [42]: plt.plot(tag_counts[0:500])
plt.title('first 500 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:500:5]), tag_counts[0:500:5])
```




```
100 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
22429 21820 20957 19758 18905 17728 15533 15097 14884 13703
13364 13157 12407 11658 11228 11162 10863 10600 10350 10224
10029 9884 9719 9411 9252 9148 9040 8617 8361 8163
8054 7867 7702 7564 7274 7151 7052 6847 6656 6553
6466 6291 6183 6093 5971 5865 5760 5577 5490 5411
5370 5283 5207 5107 5066 4983 4891 4785 4658 4549
4526 4487 4429 4335 4310 4281 4239 4228 4195 4159
4144 4088 4050 4002 3957 3929 3874 3849 3818 3797
3750 3703 3685 3658 3615 3593 3564 3521 3505 3483]
```



```
In [44]: plt.plot(tag_counts[0:100], c='b')
plt.scatter(x=list(range(0,100,5)), y=tag_counts[0:100:5], c='orange', label="quantiles with 0.25 difference")
plt.scatter(x=list(range(0,100,25)), y=tag_counts[0:100:25], c='m', label = "quantiles with 0.05 difference")

for x,y in zip(list(range(0,100,25)), tag_counts[0:100:25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500))

plt.title('first 100 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.legend()
plt.show()
print(len(tag_counts[0:100:5]), tag_counts[0:100:5])
```



```
20 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
22429 21820 20957 19758 18905 17728 15533 15097 14884 13703]
```

```
In [45]: # Store tags greater than 10K in one list
lst_tags_gt_10k = tag_df[tag_df.Counts>10000].Tags
#Print the length of the list
print ('{} Tags are used more than 10000 times'.format(len(lst_tags_gt_10k)))
# Store tags greater than 100K in one list
lst_tags_gt_100k = tag_df[tag_df.Counts>100000].Tags
#Print the length of the list.
print ('{} Tags are used more than 100000 times'.format(len(lst_tags_gt_100k)))
```

```
153 Tags are used more than 10000 times
14 Tags are used more than 100000 times
```

### Observations:

1. There are total 153 tags which are used more than 10000 times.
2. 14 tags are used more than 100000 times.
3. Most frequent tag (i.e. c#) is used 331505 times.

### 3.2.4 Tags Per Question

```
In [46]: #Storing the count of tag in each question in list 'tag_count'
tag_quest_count = tag_dtm.sum(axis=1).tolist()
#Converting each value in the 'tag_quest_count' to integer.
tag_quest_count=[int(j) for i in tag_quest_count for j in i]
print ('We have total {} datapoints.'.format(len(tag_quest_count)))


print(tag_quest_count[:5])
```

We have total 4206307 datapoints.  
[3, 4, 2, 2, 3]

```
In [47]: print( "Maximum number of tags per question: %d"%max(tag_quest_count))
print( "Minimum number of tags per question: %d"%min(tag_quest_count))
print( "Avg. number of tags per question: %f"% ((sum(tag_quest_count)*1.0)/len(tag_quest_count)))
```

Maximum number of tags per question: 5  
Minimum number of tags per question: 1  
Avg. number of tags per question: 2.899443

```
In [48]: sns.countplot(tag_quest_count, palette='gist_rainbow')
plt.title("Number of tags in the questions ")
plt.xlabel("Number of Tags")
plt.ylabel("Number of questions")
plt.show()
```



#### Observations:

1. Maximum number of tags per question: 5
2. Minimum number of tags per question: 1
3. Avg. number of tags per question: 2.899
4. Most of the questions are having 2 or 3 tags

### 3.2.5 Most Frequent Tags

```
In [ ]: # Plotting word cloud
start = datetime.now()

# Lets first convert the 'result' dictionary to 'list of tuples'
tup = dict(result.items())
#Initializing WordCloud using frequencies of tags.
wordcloud = WordCloud(    background_color='black',
                          width=1600,
                          height=800,
                          ).generate_from_frequencies(tup)


fig = plt.figure(figsize=(30,20))
plt.imshow(wordcloud)
plt.axis('off')
plt.tight_layout(pad=0)
fig.savefig("tag.png")
plt.show()
print("Time taken to run this cell :", datetime.now() - start)
```

### Observations:

A look at the word cloud shows that "c#", "java", "php", "asp.net", "javascript", "c++" are some of the most frequent tags.

### 3.2.6 The top 20 tags

```
In [50]: i=np.arange(30)
tag_df_sorted.head(30).plot(kind='bar')
plt.title('Frequency of top 20 tags')
plt.xticks(i, tag_df_sorted['Tags'])
plt.xlabel('Tags')
plt.ylabel('Counts')
plt.show()
```



**Observations:**

1. Majority of the most frequent tags are programming language.
2. C# is the top most frequent programming language.
3. Android, IOS, Linux and windows are among the top most frequent operating systems.

### 3.3 Cleaning and preprocessing of Questions

#### 3.3.1 Preprocessing

1. Sample 1M data points
2. Separate Code from Body
3. Remove Special characters from Question title and description (not in code)
4. Remove stop words (Except 'C')
5. Remove HTML Tags
6. Convert all the characters into small letters
7. Use SnowballStemmer to stem the words

```
In [51]: def striphtml(data):  
        cleanr = re.compile('<.*?>')  
        cleantext = re.sub(cleanr, ' ', str(data))  
        return cleantext  
stop_words = set(stopwords.words('english'))  
stemmer = SnowballStemmer("english")
```

```

In [52]: #http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def create_table(conn, create_table_sql):
    """ create a table from the create_table_sql statement
    :param conn: Connection object
    :param create_table_sql: a CREATE TABLE statement
    :return:
    """
    try:
        c = conn.cursor()
        c.execute(create_table_sql)
    except Error as e:
        print(e)

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the databse:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

def create_database_table(database, query):
    conn = create_connection(database)
    if conn is not None:
        create_table(conn, query)
        checkTableExists(conn)
    else:
        print("Error! cannot create the database connection.")
    conn.close()

sql_create_table = "CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text"
create_database_table("D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/Proce

```

Tables in the databse:  
QuestionsProcessed

```
In [ ]: # http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table
start = datetime.now()
read_db = 'D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train_no_dup.db'
write_db = 'D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/Processed.db'
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RAND()")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
print("Time taken to run this cell :", datetime.now() - start)
```

\_\_ we create a new data base to store the sampled and preprocessed questions \_\_

```

In [ ]: #http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table,

start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], row[2]

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=stripthtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    question=str(title)+" "+str(question)
    question=re.sub(r'^[A-Za-z]+', ' ', question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words)

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    conn_w = create_connection(write_db)
    writer =conn_w.cursor()
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post,length,is_code)")
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)

```

```
In [ ]: # dont forget to close the connections, or else you will end up with locks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()
```

```
In [ ]: if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()
```

```
In [ ]: if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()
```

```
In [ ]: #Taking 1 Million entries to a dataframe.
write_db = 'D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/Processed.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM Ques
conn_r.commit()
conn_r.close()
```

```
In [ ]: preprocessed_data.head()
```

```
In [ ]: print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

## 4. Machine Learning Models

### 4.1 Converting tags for multilable problems



X	y1	y2	y3	y4
x1	0	1	1	0
x1	1	0	0	0
x1	0	1	0	0

```
In [ ]: # binary='true' will give a binary vectorizer
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```

\_\_ We will sample the number of tags instead considering all of them (Limitation of computing power) \_\_

```
In [ ]: def tags_to_choose(n):
    t = multilabel_y.sum(axis=0).tolist()[0]
    sorted_tags_i = sorted(range(len(t)), key=lambda i: t[i], reverse=True)
    multilabel_yn=multilabel_y[:,sorted_tags_i[:n]]
    return multilabel_yn

def questions_explained_fn(n):
    multilabel_yn = tags_to_choose(n)
    x= multilabel_yn.sum(axis=1)
    return (np.count_nonzero(x==0))
```

```
In [ ]: questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/total_qs)*100))
```

```
In [ ]: fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimum is 50
print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
```

```
In [ ]: multilabel_yx = tags_to_choose(5500)
print("number of questions that are not covered :", questions_explained_fn(5500))
```

```
In [ ]: print("Number of tags in sample :", multilabel_y.shape[1])
print("number of tags taken :", multilabel_yx.shape[1], "(", (multilabel_yx.shape[1]/multilabel_y.shape[1])*100, "% of tags")
```

\_\_ We consider top 15% tags which covers 99% of the questions \_\_

## 4.2 Split the data into test and train (80:20)

```
In [ ]: total_size=preprocessed_data.shape[0]
        train_size=int(0.80*total_size)

        x_train=preprocessed_data.head(train_size)
        x_test=preprocessed_data.tail(total_size - train_size)

        y_train = multilabel_yx[0:train_size,:]
        y_test = multilabel_yx[train_size:total_size,:]
```

```
In [ ]: print("Number of data points in train data :", y_train.shape)
        print("Number of data points in test data :", y_test.shape)
```

## 4.3 Featurizing data

```
In [ ]: start = datetime.now()
        vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=True,
                                     tokenizer = lambda x: x.split(), sublinear_tf=False)
        x_train_multilabel = vectorizer.fit_transform(x_train['question'])
        x_test_multilabel = vectorizer.transform(x_test['question'])
        print("Time taken to run this cell :", datetime.now() - start)
```

```
In [ ]: print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
        print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

```
In [ ]: # https://www.analyticsvidhya.com/blog/2017/08/introduction-to-multi-label-classification/
        #https://stats.stackexchange.com/questions/117796/scikit-multi-label-classification
        # classifier = LabelPowerSet(GaussianNB())
        """

        from skmultilearn.adapt import MLkNN
        classifier = MLkNN(k=21)

        # train
        classifier.fit(x_train_multilabel, y_train)

        # predict
        predictions = classifier.predict(x_test_multilabel)
        print(accuracy_score(y_test,predictions))
        print(metrics.f1_score(y_test, predictions, average = 'macro'))
        print(metrics.f1_score(y_test, predictions, average = 'micro'))
        print(metrics.hamming_loss(y_test,predictions))

        """

        # we are getting memory error because the multilearn package is trying to convert
        # -----
        #MemoryError                                Traceback (most recent call last)
        #<ipython-input-170-f0e7c7f3e0be> in <module>()
        #----> classifier.fit(x_train_multilabel, y_train)
```

## 4.4 Applying Logistic Regression with OneVsRest

## Classifier

```
In [ ]: classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l2',
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("accuracy :",metrics.accuracy_score(y_test,predictions))
print("macro f1 score :",metrics.f1_score(y_test, predictions, average = 'macro'))
print("micro f1 scoore :",metrics.f1_score(y_test, predictions, average = 'micro'))
print("hamming loss :",metrics.hamming_loss(y_test,predictions))
print("Precision recall report :\n",metrics.classification_report(y_test, predictions))
```

```
In [ ]: from sklearn.externals import joblib
joblib.dump(classifier, 'lr_with_equal_weight.pkl')
```

## 4.5 Modeling with less data points (0.5M data points) and more weight to title

```
In [ ]: sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text, title text, body text, tags text, created_at timestamp, updated_at timestamp);
create_database_table("D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/Titlemoreweight.db", sql_create_table)
```

```
In [ ]: # http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table

read_db = 'D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/train_no_dup.db'
write_db = 'D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/Titlemoreweight.db'
train_datasize = 400000
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        # for selecting first 0.5M rows
        reader.execute("SELECT Title, Body, Tags From no_dup_train LIMIT 500001;")
        # for selecting random points
        #reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RAND() LIMIT 500001;")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
```

### 4.5.1 Preprocessing of questions

1. Separate Code from Body
2. Remove Special characters from Question title and description (not in code)

3. **Give more weightage to title : Add title three times to the question**
4. **Add 'Tags' string to the training data**
5. Remove stop words (Except 'C')
6. Remove HTML Tags
7. Convert all the characters into small letters
8. Use SnowballStemmer to stem the words

```

In [ ]: #http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table,
start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], str(row[2])

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=stripthtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    # adding title three time to the data to increase its weight
    # add tags string to the training data

    question=str(title)+" "+str(title)+" "+str(title)+" "+question

    # if questions_proccesed<=train_datasize:
    #     question=str(title)+" "+str(title)+" "+str(title)+" "+question+" "+str(tags)
    # else:
    #     question=str(title)+" "+str(title)+" "+str(title)+" "+question

    question=re.sub(r'^[A-Za-z0-9#+.\-]+', ' ', question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words)

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post,length,is_code)")
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_proccesed))

```

```
print("Time taken to run this cell :", datetime.now() - start)
```

```
In [ ]: # never forget to close the connections or else we will end up with database locks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()
```

\_\_ Sample questions after preprocessing of data \_\_

```
In [ ]: if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()
```

\_\_ Saving Preprocessed data to a Database \_\_

```
In [ ]: #Taking 1 Million entries to a dataframe.
write_db = 'D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/Titlemoreweight'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM QuestionsProcessed""")
    conn_r.commit()
    conn_r.close()
```

```
In [ ]: preprocessed_data.head()
```

```
In [ ]: print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

\_\_ Converting string Tags to multilable output variables \_\_

```
In [ ]: vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```

\_\_ Selecting 500 Tags \_\_

```
In [ ]: questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/to
```

```
In [ ]: fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimum is 50
print("with ",5500,"tags we are covering ",questions_explained[50],"% of question
```

```
In [ ]: multilabel_yx = tags_to_choose(500)
print("number of questions that are not covered :", questions_explained_fn(500),
```

```
In [ ]: import joblib

x_train = preprocessed_data.head(train_datasize)
x_test = preprocessed_data.tail(preprocessed_data.shape[0] - 400000)

y_train = multilabel_yx[0:train_datasize,:]
y_test = multilabel_yx[train_datasize:preprocessed_data.shape[0],:]

joblib.dump(x_train,"D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/x_train")
joblib.dump(x_test,"D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/x_test")
joblib.dump(y_train,"D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/y_train")
joblib.dump(y_test,"D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/y_test")
```

```
In [2]: import joblib
x_train = joblib.load("D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/x_train")
x_test = joblib.load("D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/x_test")
y_train = joblib.load("D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/y_train")
y_test = joblib.load("D:/Data Science/DataSets/Stackoverflow_Tag_Predictor/y_test")
```

```
In [3]: print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

```
Number of data points in train data : (400000, 500)
Number of data points in test data : (100000, 500)
```

```
In [4]: # Taking only 100k data points
x_train_1 = x_train
y_train_1 = y_train

x_train = x_train_1[0:80000]
x_test = x_train_1[80000:100000]
y_train = y_train_1[0:80000]
y_test = y_train_1[80000:100000]

print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (80000, 500)

Number of data points in test data : (20000, 500)

### 4.5.2 Featurizing data with Tfidf vectorizer

```
In [ ]: start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=True,
                             tokenizer = lambda x: x.split(), sublinear_tf=False)
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

```
In [ ]: print("Dimensions of train data X:", x_train_multilabel.shape, "Y :", y_train.shape)
print("Dimensions of test data X:", x_test_multilabel.shape, "Y:", y_test.shape)
```

### 4.5.3 Applying Logistic Regression with OneVsRest Classifier



```
In [ ]: start = datetime.now()
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l2'))
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)
```

```
In [ ]: joblib.dump(classifier, 'lr_with_more_title_weight.pkl')
```

```
In [ ]: start = datetime.now()
classifier_2 = OneVsRestClassifier(LogisticRegression(penalty='l1'), n_jobs=-1)
classifier_2.fit(x_train_multilabel, y_train)
predictions_2 = classifier_2.predict(x_test_multilabel)
print("Accuracy :",metrics.accuracy_score(y_test, predictions_2))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions_2))
print("Time taken to run this cell :", datetime.now() - start)
```

## 5. Assignments

1. Use bag of words upto 4 grams and compute the micro f1 score with Logistic regression(OvR)
2. Use tldif vectorizer upto 4 grams and compute the micro f1 score with Knearest (OvR)
3. Add some extra features and try to get micro f1 score > 0.5

### Featurizing data with BoW vectorizer

```
In [5]: start = datetime.now()
print("Starting time: ",start)
vectorizer = CountVectorizer(max_features=500,tokenizer = lambda x: x.split(), n
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

```
Starting time: 2019-04-12 18:01:26.979977
Time taken to run this cell : 0:04:43.896785
```

```
In [6]: print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shap
print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

```
Dimensions of train data X: (80000, 500) Y : (80000, 500)
Dimensions of test data X: (20000, 500) Y: (20000, 500)
```

### Logistic Regression with BoW

```

In [28]: #Bow
start = datetime.now()
print("Starting time: ",start)

classifier_2 = OneVsRestClassifier(LogisticRegression(penalty='l1'))
classifier_2.fit(x_train_multilabel, y_train)
predictions_2 = classifier_2.predict(x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions_2))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

print (metrics.classification_report(y_test, predictions_2))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Starting time: 2019-04-10 20:25:10.380943
Accuracy : 0.13055
Hamming loss 0.0036048
Micro-average quality numbers
Precision: 0.5372, Recall: 0.1861, F1-measure: 0.2765

```

```

In [7]: print(np.shape(x_train_multilabel))
print(np.shape(y_train))

```

```

(80000, 500)
(80000, 500)

```

```
In [8]: from sklearn.model_selection import GridSearchCV
import datetime

#hyper-paramater tuning
startTime3 = datetime.datetime.now()
print("Current Time = ",startTime3)

#Using GridSearchCV with L1 Regularizer
tuned_parameters = [{'estimator__C': [100,0.1, 0.0001]}]

classifier = OneVsRestClassifier(LogisticRegression(penalty='l1'))

model_l2 = GridSearchCV(classifier,tuned_parameters,scoring = 'f1_micro', cv=3)
model_l2.fit(x_train_multilabel, y_train)
```

C:\Users\AbhiShek\Anaconda3\lib\site-packages\sklearn\multiclass.py:76: UserWarning: Label not 429 is present in all training examples.  
 str(classes[c]))  
C:\Users\AbhiShek\Anaconda3\lib\site-packages\sklearn\multiclass.py:76: UserWarning: Label not 432 is present in all training examples.  
 str(classes[c]))  
C:\Users\AbhiShek\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: ConvergenceWarning: LibSVM failed to converge. Increase the number of iterations

```
In [10]: GS_OPTIMAL_clf_sgd = model_l2.best_estimator_
print("GS_OPTIMAL_clf_sgd = ",GS_OPTIMAL_clf_sgd)
best_score_model_l2 = model_l2.best_score_
print("\nBest score: ",best_score_model_l2)
test_score_l2 = model_l2.score(x_test_multilabel, y_test)
print("test_score_l2 = ",test_score_l2)
param_C = model_l2.best_params_["estimator__C"]
print("Best C= ",param_C)
```

```
GS_OPTIMAL_clf_sgd = OneVsRestClassifier(estimator=LogisticRegression(C=0.1, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l1', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False),
    n_jobs=None)
```

```
Best score: 0.34329031831081575
test_score_l2 = 0.25842721112152184
Best C= 0.1
```

## Applying Logistic regression with best C value

```
In [17]: #BoW
start = datetime.now()
print("Starting time: ",start)

classifier_2 = OneVsRestClassifier(LogisticRegression(penalty='l1',C=0.1,random_
classifier_2.fit(x_train_multilabel, y_train)
predictions_2 = classifier_2.predict(x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions_2))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

print (metrics.classification_report(y_test, predictions_2))
print("Time taken to run this cell :", datetime.now() - start)
```

Starting time: 2019-04-13 14:44:03.450358

Accuracy : 0.13855

Hamming loss 0.0034033

Micro-average quality numbers

Precision: 0.6669, Recall: 0.1603, F1-measure: 0.2584

Macro-average quality numbers

Precision: 0.1743, Recall: 0.0493, F1-measure: 0.0706

	precision	recall	f1-score	support
0	0.92	0.47	0.62	1424
1	0.64	0.08	0.14	1931
2	0.82	0.32	0.46	1374
3	0.83	0.36	0.50	797
4	0.84	0.41	0.55	1311
5	0.80	0.27	0.41	820
6	0.90	0.45	0.60	1014
7	0.87	0.52	0.65	702
8	0.60	0.12	0.21	544
9	0.83	0.54	0.65	269
10	0.50	0.10	0.20	100

### OneVsRestClassifier with Linear SVM with BoW

```
In [34]: #Bow
start = datetime.now()
print("Starting time: ",start)

classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', alpha=0.00001, penalti
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict (x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)
```

Starting time: 2019-04-11 11:36:23.508911

Accuracy : 0.0671

Hamming loss 0.0074196

Micro-average quality numbers

Precision: 0.1675, Recall: 0.2532, F1-measure: 0.2016

Macro-average quality numbers

Precision: 0.0701, Recall: 0.1155, F1-measure: 0.0785

C:\Users\AbhiShek\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1145: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples.

'recall', 'true', average, warn\_for)

C:\Users\AbhiShek\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1145: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no true samples.

'recall', 'true', average, warn\_for)

C:\Users\AbhiShek\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1145: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples.

'recall', 'true', average, warn\_for)

## Featurizing data with tfidf vectorizer

```
In [32]: start = datetime.now()
print("Starting time: ",start)
vectorizer_tfidf = TfidfVectorizer(min_df=0.00009, max_features=500, smooth_idf=
tokenizer = lambda x: x.split(), sublinear_tf=False)
x_train_multilabel_tfidf = vectorizer_tfidf.fit_transform(x_train['question'])
x_test_multilabel_tfidf = vectorizer_tfidf.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Starting time: 2019-04-11 11:30:14.504213  
Time taken to run this cell : 0:05:11.062140

```
In [33]: print("Dimensions of train data X:",x_train_multilabel_tfidf.shape, "Y :",y_train_multilabel_tfidf.shape)
print("Dimensions of test data X:",x_test_multilabel_tfidf.shape,"Y:",y_test_multilabel_tfidf.shape)
```

Dimensions of train data X: (80000, 500) Y : (80000, 500)  
Dimensions of test data X: (20000, 500) Y: (20000, 500)

```
In [20]: #tfidf
start = datetime.now()
print("Starting time: ",start)

classifier_2 = OneVsRestClassifier(LogisticRegression(penalty='l1'), n_jobs=-1)
classifier_2.fit(x_train_multilabel_tfidf, y_train)
predictions_2 = classifier_2.predict(x_test_multilabel_tfidf)

print("Accuracy :",metrics.accuracy_score(y_test, predictions_2))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

print (metrics.classification_report(y_test, predictions_2))
print("Time taken to run this cell :", datetime.now() - start)
```

```
Starting time: 2019-04-10 18:13:28.793510
Accuracy : 0.1435
Hamming loss 0.0033437
Micro-average quality numbers
Precision: 0.7030, Recall: 0.1668, F1-measure: 0.2696
Macro-average quality numbers
Precision: 0.1871, Recall: 0.0537, F1-measure: 0.0753
```



```

In [26]: #tfidf
start = datetime.now()
print("Starting time: ",start)

classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', alpha=0.00001, penalti
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict (x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,

print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

Starting time: 2019-04-10 19:59:50.760367

Accuracy : 0.14565

Hamming loss 0.0032959

Micro-average quality numbers

Precision: 0.7996, Recall: 0.1458, F1-measure: 0.2466

Macro-average quality numbers

Precision: 0.1008, Recall: 0.0355, F1-measure: 0.0468

	precision	recall	f1-score	support
0	0.94	0.51	0.67	1424
1	0.77	0.01	0.01	1931
2	0.85	0.31	0.46	1374
3	0.82	0.34	0.48	797
4	0.88	0.35	0.50	1311
5	0.80	0.32	0.46	820
6	0.89	0.45	0.60	1014
7	0.88	0.55	0.68	702
8	0.61	0.13	0.22	544
9	0.84	0.51	0.63	269
10	0.00	0.00	0.00	105

```
In [2]: from prettytable import PrettyTable
x = PrettyTable()

x.field_names = ["Vectorizer", "Model", "n_gram range", "F1_Score"]

x.add_row(["BoW", "Logistic Regression", 4, 0.2765])
x.add_row(["BoW", "Log. Regression (Grid Search)", 4, 0.2584])
x.add_row(["BoW", "Linear SVM", 4, 0.2016])

x.add_row(["TFIDF", "Logistic Regression", 4, 0.2696])
x.add_row(["TFIDF", "Lr. SVM", 4, 0.2466])

print(x)
```

Vectorizer	Model	n_gram range	F1_Score
BoW	Logistic Regression	4	0.2765
BoW	Log. Regression (Grid Search)	4	0.2584
BoW	Linear SVM	4	0.2016
TFIDF	Logistic Regression	4	0.2696
TFIDF	Lr. SVM	4	0.2466

In [ ]: